

ANOVA

Team 14

Eleanor LaRocco, Erik Zwicklbauer, Junghee Mun, Sebastian Rincon

What is ANOVA?

- **ANOVA: ANalysis Of VAriance**
 - Analyzing variance to look for differences
- Testing how well nested models fit a dataset in comparing against models that add features
 - Less complex models (M_0) vs More complex models (M_1)
- R: `anova()` function performs a hypothesis test comparing two models
- Null hypothesis (H_0): $M_0 = M_1$
 - Model 0 and Model 1 explains the data equally well
 - p-value is the probability that the null is true given the data
 - If the p-value is very small (ie. 0.05), reject the H_A and accept H_0
- Alternative hypothesis (H_A): $M_0 \neq M_1$

ANOVA's Assumptions

- Models must be **nested**
 - The predictors in M_0 must be a subset of the predictors in M_1
- **Equal variances** between models (homoscedasticity)

Interpreting ANOVA

Hypothesis	Correct model?	R formula for correct model
Null	M0	$Y \sim A + B$
Alternative	M1	$Y \sim A + B + C$

- Null: $\hat{y}_{\text{wage}} = \beta_1 \text{age} + \beta_2 \text{age}^2$
- Alternative: $\hat{y}_{\text{wage}} = \beta_1 \text{age} + \beta_2 \text{age}^2 + \beta_3 \text{age}^3$

Contents of ANOVA (1)

- **Res. Df:** Residual Degrees of Freedom
- **RSS:** Residual Sum of Squares

$$RSS = \sum_{i=1}^n (y_i - f(x_i))^2$$

- **Df:** Degrees of Freedom
- **SS:** Sum of Squares
 - $SS \text{ model}_1 = RSS \text{ model}_0 - RSS \text{ model}_1$

Analysis of Variance Table

Model 1	1: wage ~ poly(age, 1)						
Model 2	2: wage ~ poly(age, 2)						
Model 3	3: wage ~ poly(age, 3)						
Model 4	4: wage ~ poly(age, 4)						
Model 5	5: wage ~ poly(age, 5)						
Model 6	6: wage ~ poly(age, 6)						
Model 7	7: wage ~ poly(age, 7)						
Model 8	8: wage ~ poly(age, 8)						
Model 9	9: wage ~ poly(age, 9)						
Model 10	10: wage ~ poly(age, 10)						
	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)	
1	2998	5022216					
2	2997	4793430	1	228786	143.7638	< 2.2e-16	***
3	2996	4777674	1	15756	9.9005	0.001669	**
4	2995	4771604	1	6070	3.8143	0.050909	.
5	2994	4770322	1	1283	0.8059	0.369398	
6	2993	4766389	1	3932	2.4709	0.116074	
7	2992	4763834	1	2555	1.6057	0.205199	
8	2991	4763707	1	127	0.0796	0.777865	
9	2990	4756703	1	7004	4.4014	0.035994	*
10	2989	4756701	1	3	0.0017	0.967529	

Contents of ANOVA (2)

- **F-statistic:**

- `SS.model1 <- RSS.model0 - RSS.model1`
- `res <- RSS.model1/Res.Df.model1`
- `diff <- SS.model1/(Res.Df.model0-Res.Df.model1)`
- `F.stat <- diff/res`

- **P-value:** Probability that the null is true given the data

- P-value > 0.05: Model 1 and 2 are not statistically different
- P-value < 0.05: Model 1 is an improvement from Model 0

Analysis of Variance Table

Model 1: wage ~ poly(age, 1)
Model 2: wage ~ poly(age, 2)
Model 3: wage ~ poly(age, 3)
Model 4: wage ~ poly(age, 4)
Model 5: wage ~ poly(age, 5)
Model 6: wage ~ poly(age, 6)
Model 7: wage ~ poly(age, 7)
Model 8: wage ~ poly(age, 8)
Model 9: wage ~ poly(age, 9)
Model 10: wage ~ poly(age, 10)

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)	
1	2998	5022216					
2	2997	4793430	1	228786	143.7638	< 2.2e-16	***
3	2996	4777674	1	15756	9.9005	0.001669	**
4	2995	4771604	1	6070	3.8143	0.050909	.
5	2994	4770322	1	1283	0.8059	0.369398	
6	2993	4766389	1	3932	2.4709	0.116074	
7	2992	4763834	1	2555	1.6057	0.205199	
8	2991	4763707	1	127	0.0796	0.777865	
9	2990	4756703	1	7004	4.4014	0.035994	*
10	2989	4756701	1	3	0.0017	0.967529	

Check your ANOVA knowledge!

